



- **Label:** It indicates the nature of the report, which can be one of the following: *bug*, *enhancement*, and *question*.
- **Issue title:** A short descriptive sentence that indicates at a glance what the issue is about.
- **Issue body:** Each report must contain an issue body that includes a description explaining the purpose of the issue. This can consist of any details that might help resolve the problem.
- **Issue URL:** the URL to access the report on GitHub.
- **Repository URL:** Every issue is associated with a GitHub repository. This metadata stores the link of the remote GitHub repository.
- **Creation timestamp:** the timestamp of when the report was created.
- **Author association:** It describes how the issue creator is related to the repository. There are five types of author associations: *Owner*, *Collaborator*, *Contributor*, *Member*, *None*.

The dataset is split into a *training set* and a *test set* by the organizers of the NLBSE'22 tool competition [7]. The training set contains **722,899** (90%) and the test set contains **80,518** (10%) of the total labeled issue reports. We split the training set into a new train and validation set. The new train set contains 85% of the previous train set, and 15% of the data is in the validation set. We used `train_test_split` function from *scikit-learn* [14] where we re-shuffled the data and split in a stratified fashion, using the class labels because our dataset is imbalanced. The dataset distribution is given in Table 1.

Set	Bug	Enhancement	Question	Total
Training set	306,937 (50.0%)	254,468 (41.4%)	53,059 (8.6%)	614,464
Validation set	54,166 (50.0%)	44,906 (41.4%)	9,363 (8.6%)	108,435
Testing set	40,152 (49.9%)	33,290 (41.3%)	7,076 (8.8%)	80,518

Table 1: Dataset distribution.

## 2.2 Preprocessing

Before training the model, we preprocessed the data and fixed the model's hyperparameters. This section discusses data preprocessing steps and details of the hyper-parameters (i.e., Pre-trained Model, Optimizer, and Scheduler).

**2.2.1 Text Cleaning & Feature Extraction.** A GitHub issue contains a *title* and a detailed *description*. An issue's description may contain code segments or screenshots of the output. For example, Figure 2 shows a GitHub issue from Facebook's Flow repository containing code segments. Hence, we need to clean the data before encoding it.

As previously explained in Section 2.1, each issue has seven metadata in which one of them is the *label*. We consider the issue *title* and the issue *body* as the main features out of six features.

First, we concatenate these two features (*title* and *body*) into a new metadata, which we refer to as *issue data*. Then, we used the *Gensim*<sup>2</sup> library to remove repeating whitespace characters (i.e., spaces, tabs, and line breaks) from the *issue data*. Moreover, we replaced tabs and line breaks with spaces. This processed issue data is used as the feature of our model.

<sup>2</sup><https://radimrehurek.com/gensim/>

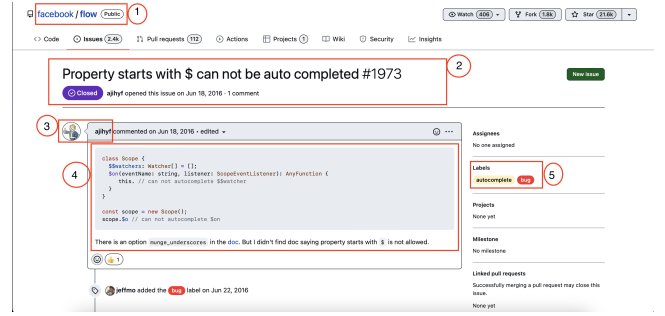


Figure 2: GitHub issue with bug tag from Facebook/flow  
1. Repository Name, 2. Issue title, 3. Author handle and avatar, 4. Issue body, 5. Label

**2.2.2 Encoding.** To use a pre-trained BERT model, the feature data must be divided into *tokens*, then mapped to their respective indexes in the tokenizer vocabulary. We used *BertTokenizer*<sup>3</sup> implementation, which is available in Hugging Face's transformers [18] package. We used the 'bert-base-uncased' pre-trained model [4] to tokenize our issue data. The 'bert-base-uncased' model is trained in English using a masked language modeling (MLM) objective. This model is uncased (i.e., it is case-insensitive). For instance, it does not differentiate between "issue" and "Issue".

BERT is a pre-trained model that requires input data in a specific format. Thus, we need the following items [11]:

- **[SEP]:** This token marks the end of a sentence or the separation between two sentences
- **[CLS]:** This token is used at the beginning of our text and is used for classification tasks, but BERT expects it no matter what your application is.
- **Token:** It complies with the fixed vocabulary used in BERT.
- **Token ID:** It is for the token generated from BERT's tokenizer.
- **Mask ID:** It indicates which elements in the sequence are tokens and which are padding elements.
- **Segment IDs:** It distinguishes different sentences.
- **Positional Embedding:** It shows token position within the sequence.

We used the `batch_encode_plus` method from *BertTokenizer* to handle the specific format described before. We used this method to extract the attention mask with the following parameters:

- `add_special_token` is set to **true** to encode special token;
- `return_attention_mask` is set to **true** to get the attention mask;
- `padding` is set to **'longest'** to pad to the longest sequence in the batch;
- `truncation` is set to **true** to truncate to the maximum acceptable input length for the model;
- `return_tensors` is set to **'pt'** to get PyTorch tensors as a return value.

We used *token ids* (`input_ids`) and *attention mask* to create a *TensorDataset* which is later fed into *DataLoader* to train and evaluate the model. The data loader is configured with random sampling, and batch size equals four. Although the authors of the

<sup>3</sup>[https://huggingface.co/docs/transformers/model\\_doc/bert#transformers.BertTokenizer](https://huggingface.co/docs/transformers/model_doc/bert#transformers.BertTokenizer)

BERT paper recommended using a batch size equal to 16 or 32 [4], our GPU has a memory limitation, causing out-of-memory errors when using the recommended sizes. Therefore, we used a shorter batch size to solve this memory constraint at the cost of increasing the training time.

**2.2.3 Pretrained Model.** We used the BERT pre-trained model from Google AI [4]. The use of the bidirectional training of Transformer [17] for language modeling is BERT’s fundamental technological breakthrough. This contrasts with previous efforts, which looked at a text sequence from left to right or combined left-to-right and right-to-left training. [12]. The research findings suggest that bidirectionally trained language models can better understand language context and flow than single-direction language models.

We first alter the pre-trained BERT model to provide classification outputs in our work. Then, we keep training the model on our dataset until the complete model, end-to-end, is well-suited to our objective. We used BertForSequenceClassification<sup>4</sup> from Hugging Face. This is the standard BERT model with a single classification layer placed on top, which we employed as a document classifier. The pre-trained BERT model and the additional untrained classification layer get trained on our dataset. We utilized the pre-trained model "bert-base-uncased," which refers to the version with just lowercase characters ("uncased"). Since we did not want the model to return the attention weights and all hidden states, we disabled the flags while initializing the pre-trained model. After initializing the model, we fix our optimizer and scheduler, as described in the next subsections.

**2.2.4 Optimizer.** We used AdamW [10] optimizer for the training. We used the implementation of the Adam algorithm with weight decay fix from HuggingFace. We used  $1e^{-5}$  as the learning rate (lr) and  $1e^{-8}$  as the eps parameter, which is a very small number to prevent any division by zero in the implementation.

**2.2.5 Learning Rate Schedule.** The learning rate schedule is a hyperparameter that changes the learning rate between epochs or iterations to minimize the model’s loss. We used a *linear schedule with warmup* implementation from HuggingFace. It created a schedule with a learning rate that decreases linearly from the initial lr set in the optimizer to 0, after a warmup period during which it increases linearly from 0 to the initial lr fixed in the optimizer. The implement takes input for the number of warmup steps which is 0 and the number of training steps, which is the size of the data loader of the train set multiplying with the iteration number.

## 2.3 Training

We used a previously created data loader to unpack the batch in our training phase, and each tensor was copied to the GPU. After clearing any previously calculated gradients, we performed a forward pass. In this step, the model provided the loss and logits as the output before the activation. Then, we perform a backward pass for calculating the gradients. We also clipped the norm of the gradients to 1.0. This is to help prevent the "exploding gradients" problem [15]. Then, we perform the optimizer’s step to dictate how the parameters are modified based on their gradients, the learning

rate, etc. Finally, we performed the scheduler’s step to update the learning rate.

After each iteration, we calculated the average training loss and model performance on the validation set by calculating validation loss. We saved the model states for evaluating the model on training sets. Since the authors of BERT recommended having 2-4 iterations, we had **four iterations** for training.

## 2.4 Evaluation

We evaluate our model with the following metrics:

- **Precision (P):** It is calculated by dividing the number of records with correctly predicted labels by the total number of predicted observations in that class:  $P = \frac{TP}{TP+FP}$ . Here, TP (true positives) is the number of records for which the label is predicted correctly. In contrast, FP (false positives) denotes the number of records for which the label is incorrectly predicted.
- **Recall (R):** It is computed for each group A by dividing the number of successfully predicted observations in A by the total number of observations in the corresponding class:  $R = \frac{TP}{TP+FN}$ . Here, FN (false negatives) is the number of observations in class A which are falsely predicted as other labels.
- **F1-Score (F1):** The F1-score combines the precision and recall of a classifier into a single metric by taking their harmonic mean:  $F1 = \frac{2 \times (P \times R)}{P + R}$ .

We calculated the Precision, Recall, and F1-Score for each class. We used micro-averaging as the cross-class aggregation method to calculate global scores due to the class imbalance present in the data.

## 2.5 Implementation Details

We run our training job on a computing node that consists of Dual Twelve-core 2.2GHz Intel Xeon processors - 24 total cores, 128 GB RAM, and 4 NVIDIA GeForce GTX 1080 Ti GPU accelerators. We used a single core and one GPU to train our model. It takes approximately 18 hours to train for a single iteration. We used PyTorch [13] as a Deep learning framework, as well as the tokenizer, pre-trained model, optimizer, and scheduler from HuggingFace<sup>5</sup>. We used the evaluation matrices implementation from Scikit-learn [14].

## 3 RESULTS

We used BERT pre-trained model to tune our dataset for classifying GitHub issue report into three classes. We compare our result with a baseline approach that used FastText [2]. Table 2 summarizes our results. We used micro-averaging as the cross-class aggregation method to calculate global scores. For that reason, the Precision (P), Recall (R), and F1-Score (F1) have the same value, where the highest value is **0.8586**, and the average is **0.8571**.

Our model consistently outperformed the baseline model using FastText. It also achieved a better result in every class regarding the F1 score. It especially outperformed the baseline approach with respect to classifying issues that are labeled as *questions*.

<sup>4</sup>[https://huggingface.co/docs/transformers/model\\_doc/bert#transformers.BertForSequenceClassification](https://huggingface.co/docs/transformers/model_doc/bert#transformers.BertForSequenceClassification)

<sup>5</sup><https://huggingface.co/>

Model	Metrics	Bug	Enhancement	Question	Global
FastText	Precision	0.8314	0.8155	0.6521	0.8162
	Recall	0.8725	0.8464	0.3502	
	F1-Score	0.8515	0.8307	0.4557	
Epoch 1	Precision	0.8755	0.8486	0.7269	0.8554
	Recall	0.8913	0.8934	0.4708	
	F1-Score	0.8833	0.8704	0.5714	
Epoch 2	Precision	0.8831	0.8591	0.6789	<b>0.8586</b>
	Recall	0.8886	0.8880	0.5469	
	F1-Score	0.8859	<b>0.8733</b>	<b>0.6058</b>	
Epoch 3	Precision	0.8722	0.8573	0.7392	0.8584
	Recall	0.90146	0.8875	0.4740	
	F1-Score	<b>0.8866</b>	0.8721	0.5776	
Epoch 4	Precision	0.8763	0.8631	0.6706	0.8561
	Recall	0.8927	0.8772	0.5466	
	F1-Score	0.8844	0.8701	0.6023	

**Table 2: Result Comparison Between FastText and our BERT-based Model.**

## 4 RELATED WORK

Kallis *et al.* [9] used FastText [2] to predict the types of GitHub issues by using issue titles and description as features. They built Ticket Tagger, a GitHub app to help developers in assigning issue types [8]. They achieved 0.75, 0.74, and 0.48 F1-score for bug, enhancement, and question, respectively, by training on a balanced set and testing an unbalanced set.

Artmann *et al.* [1] investigated the use of linear regression (LR), convolutional neural network (CNN), recurrent neural network (RNN), random forest (RF), and k-nearest-neighbor (KNN) algorithms for a multi-label text classification of GitHub issue reports. They used a 38,000 training rows dataset, and a test set containing around 12,000 rows. They split their dataset into three smaller datasets with different labels. The CNN algorithm achieved the highest F1-score for every data set.

Fan *et al.* [6] studied text-based classification approaches on a large-scale dataset of GitHub issue reports. Four different machine learning classifiers (i.e., Support Vector Machine - SVM, Naive Bayes, Logistic Regression, and Random Forest) were evaluated using 80 popular projects in GitHub consisting of about 252,000 issues. They labeled the issues into two classes: *bugs* and *non-bugs*. They introduced a new matrix based on F1-score, average F-measure as  $f_{avg}$ , F-measure of bug (nonbug) as  $f_{bug}$  ( $f_{nonbug}$ ), and number of bug (nonbug) as  $n_{bug}$  ( $n_{nonbug}$ ).

$$f_{avg} = \frac{n_{bug} * f_{bug} + n_{nonbug} * f_{nonbug}}{n_{bug} + n_{nonbug}} \quad (1)$$

They observed that text-based classification approaches can achieve 69.7% to 98.9% of average F-measure (calculated as Equation 1) on their dataset. They also found that the SVM classifier was the most effective approach compared to other typical classifiers.

## 5 CONCLUSION

Automated issue type classification can be very helpful during software maintenance, specially in open source projects where many users can open issues. This paper discussed a BERT-Based approach to automatically label issues as a question, bug, or enhancement.

Our model achieved an F1-Score of 0.8586 (on average), indicating that it can be used to predict issue report class to reduce manual work. In the future, we aim to improve our approach with a larger dataset, especially for issue reports with question tags. This approach can be integrated as a GitHub extension.

## REFERENCES

- [1] Daniel Artmann. 2020. Applying machine learning algorithms to multi-label text classification on GitHub issues.
- [2] Piotr Bojanowski, Edouard Grave, Armand Joulin, and Tomas Mikolov. 2017. Enriching Word Vectors with Subword Information. arXiv:1607.04606 [cs.CL]
- [3] Javier Luis Cánovas Izquierdo, Valerio Cosentino, Belén Rolandi, Alexandre Bergel, and Jordi Cabot. 2015. GiLA: GitHub label analyzer. In *2015 IEEE 22nd International Conference on Software Analysis, Evolution, and Reengineering (SANER)*. 479–483. <https://doi.org/10.1109/SANER.2015.7081860>
- [4] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2018. BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding. CoRR abs/1810.04805 (2018). arXiv:1810.04805 <http://arxiv.org/abs/1810.04805>
- [5] Andrea Di Sorbo, Giovanni Grano, Corrado Aaron Visaggio, and Sebastiano Panichella. 2021. Investigating the criticality of user-reported issues through their relations with app rating. *Journal of Software: Evolution and Process* 33, 3 (2021), e2316.
- [6] Qiang Fan, Yue Yu, Gang Yin, Tao Wang, and Huaimin Wang. 2017. Where Is the Road for Issue Reports Classification Based on Text Mining?. In *2017 ACM/IEEE International Symposium on Empirical Software Engineering and Measurement (ESEM)*. 121–130. <https://doi.org/10.1109/ESEM.2017.19>
- [7] Rafael Kallis, Oscar Chaparro, Andrea Di Sorbo, and Sebastiano Panichella. 2022. NLBSE'22 Tool Competition. In *Proceedings of The 1st International Workshop on Natural Language-based Software Engineering (NLBSE'22)*.
- [8] Rafael Kallis, Andrea Di Sorbo, Gerardo Canfora, and Sebastiano Panichella. 2019. Ticket Tagger: Machine Learning Driven Issue Classification. In *2019 IEEE International Conference on Software Maintenance and Evolution (ICSME)*. 406–409. <https://doi.org/10.1109/ICSME.2019.00070>
- [9] Rafael Kallis, Andrea Di Sorbo, Gerardo Canfora, and Sebastiano Panichella. 2021. Predicting issue types on GitHub. *Science of Computer Programming* 205 (2021), 102598. <https://doi.org/10.1016/j.scico.2020.102598>
- [10] Ilya Loshchilov and Frank Hutter. 2019. Decoupled Weight Decay Regularization. arXiv:1711.05101 [cs.LG]
- [11] Chris McCormick and Nick Ryan. 2019. Bert word embeddings tutorial. <https://mccormickml.com/2019/05/14/BERT-word-embeddings-tutorial/#2-input-formatting>
- [12] Oren Melamud, Jacob Goldberger, and Ido Dagan. 2016. context2vec: Learning Generic Context Embedding with Bidirectional LSTM. In *Proceedings of The 20th SIGNLL Conference on Computational Natural Language Learning*. Association for Computational Linguistics, Berlin, Germany, 51–61. <https://doi.org/10.18653/v1/K16-1006>
- [13] Adam Paszke, Sam Gross, Francisco Massa, et al. 2019. PyTorch: An Imperative Style, High-Performance Deep Learning Library. In *Advances in Neural Information Processing Systems* 32, H. Wallach, H. Larochelle, A. Beygelzimer, F. d'Alché-Buc, E. Fox, and R. Garnett (Eds.). Curran Associates, Inc., 8024–8035. <http://papers.neurips.cc/paper/9015-pytorch-an-imperative-style-high-performance-deep-learning-library.pdf>
- [14] Fabian Pedregosa, Gaël Varoquaux, Alexandre Gramfort, et al. 2018. Scikit-learn: Machine Learning in Python. arXiv:1201.0490 [cs.LG]
- [15] George Philipp, Dawn Song, and Jaime G. Carbonell. 2018. The exploding gradient problem demystified - definition, prevalence, impact, origin, tradeoffs, and solutions. arXiv:1712.05577 [cs.LG]
- [16] S Shylesh. 2017. A study of software development life cycle process models. In *National Conference on Reinventing Opportunities in Management, IT, and Social Sciences*. 534–541.
- [17] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, and Illia Polosukhin. 2017. Attention Is All You Need. arXiv:1706.03762 [cs.CL]
- [18] Thomas Wolf, Lysandre Debut, Victor Sanh, Julien Chaumond, et al. 2020. Transformers: State-of-the-Art Natural Language Processing. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing: System Demonstrations*. Association for Computational Linguistics, Online, 38–45. <https://www.aclweb.org/anthology/2020.emnlp-demos.6>
- [19] Thomas Zimmermann, R. Premraj, Jonathan Sillito, and Silvia Breu. 2009. Improving Bug Tracking Systems. 247 – 250. <https://doi.org/10.1109/ICSE-COMPANION.2009.5070993>